



Introduction to Data Mining

Lecture #22: Dimensionality Reduction-2

U Kang
Seoul National University



In This Lecture

- Learn how to answer queries using SVD
- Learn the motivation and definition of CUR decomposition, an alternative method for SVD
- Compare CUR and SVD



Outline

- ➔ ☐ SVD Case Studies
- ☐ CUR Decomposition



Case study: How to query?

- Q: Find users that like 'Matrix'
- A: Map query into a 'concept space' – how?

Diagram illustrating the mapping of a query into a concept space:

Query Matrix (SciFi vs Romance):

	Matrix	Alien	Serenity	Casablanca	Amelie
SciFi	1	1	1	0	0
	3	3	3	0	0
	4	4	4	0	0
	5	5	5	0	0
Romance	0	2	0	4	4
	0	0	0	5	5
	0	1	0	2	2

Concept Space Matrix:

0.13	0.02
0.41	0.07
0.55	0.09
0.68	0.11
0.15	-0.59
0.07	-0.73
0.07	-0.29

Transformation Matrix:

12.4	0
0	9.5

Resulting Matrix:

0.56	0.59	0.56	0.09	0.09
0.12	-0.02	0.12	-0.69	-0.69

U Kang

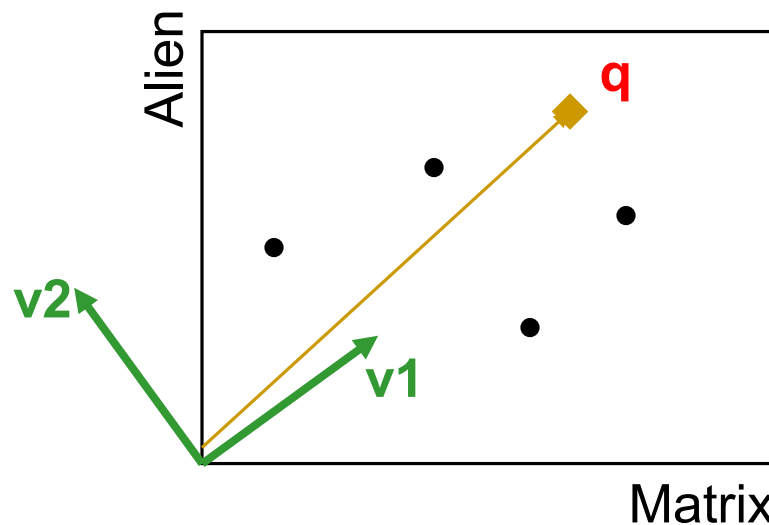


Case study: How to query?

- **Q:** Find users that like 'Matrix'
- **A:** Map query into a 'concept space' – how?

$$\mathbf{q} = \begin{matrix} & \text{Matrix} & \text{Alien} & \text{Serenity} & \text{Casablanca} & \text{Amelie} \\ \begin{bmatrix} 5 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

Project into concept space:
Inner product with each
'concept' vector \mathbf{v}_i



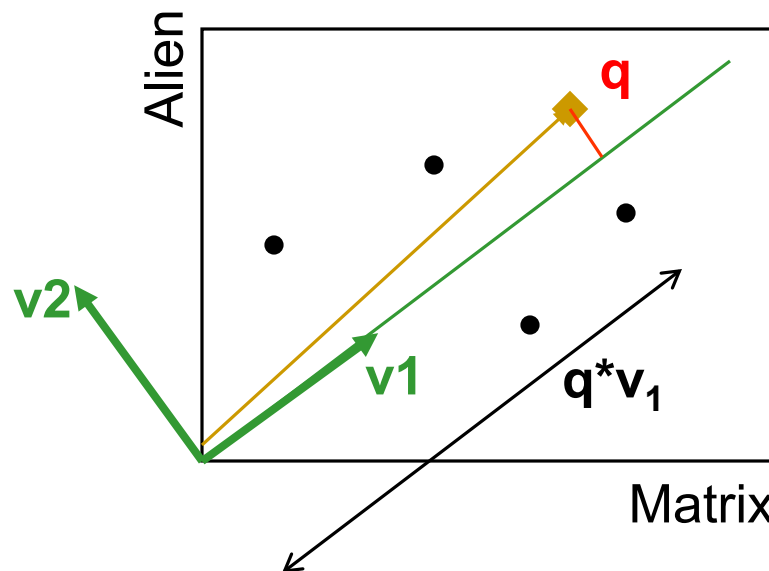


Case study: How to query?

- Q: Find users that like 'Matrix'
- A: Map query into a 'concept space' – how?

$$\mathbf{q} = \begin{matrix} & \text{Matrix} & \text{Alien} & \text{Serenity} & \text{Casablanca} & \text{Amelie} \\ \begin{bmatrix} 5 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

Project into concept space:
Inner product with each
'concept' vector \mathbf{v}_i





Case study: How to query?

Compactly, we have:

$$\mathbf{q}_{\text{concept}} = \mathbf{q} \mathbf{V}$$

$$\mathbf{q} = \begin{matrix} & \text{Matrix} & \text{Alien} & \text{Serenity} & \text{Casablanca} & \text{Amelie} \\ \begin{bmatrix} 5 & 0 & 0 & 0 & 0 \end{bmatrix} & \times & \begin{bmatrix} 0.56 & 0.12 \\ 0.59 & -0.02 \\ 0.56 & 0.12 \\ 0.09 & -0.69 \\ 0.09 & -0.69 \end{bmatrix} & = & \begin{bmatrix} 2.8 & 0.6 \end{bmatrix} \end{matrix}$$

movie-to-concept similarities (V)

SciFi-concept



Case study: How to query?

- How would the user d that rated ('Alien', 'Serenity') be handled?

$$\mathbf{d}_{\text{concept}} = \mathbf{d} \mathbf{V}$$

$$\mathbf{q} = \begin{bmatrix} \text{Matrix} \\ 0 \\ \text{Alien} \\ 4 \\ \text{Serenity} \\ 5 \\ \text{Casablanca} \\ 0 \\ \text{Amelie} \\ 0 \end{bmatrix} \mathbf{x} \begin{bmatrix} 0.56 & 0.12 \\ 0.59 & -0.02 \\ 0.56 & 0.12 \\ 0.09 & -0.69 \\ 0.09 & -0.69 \end{bmatrix} = \begin{bmatrix} 5.2 & 0.4 \end{bmatrix}$$

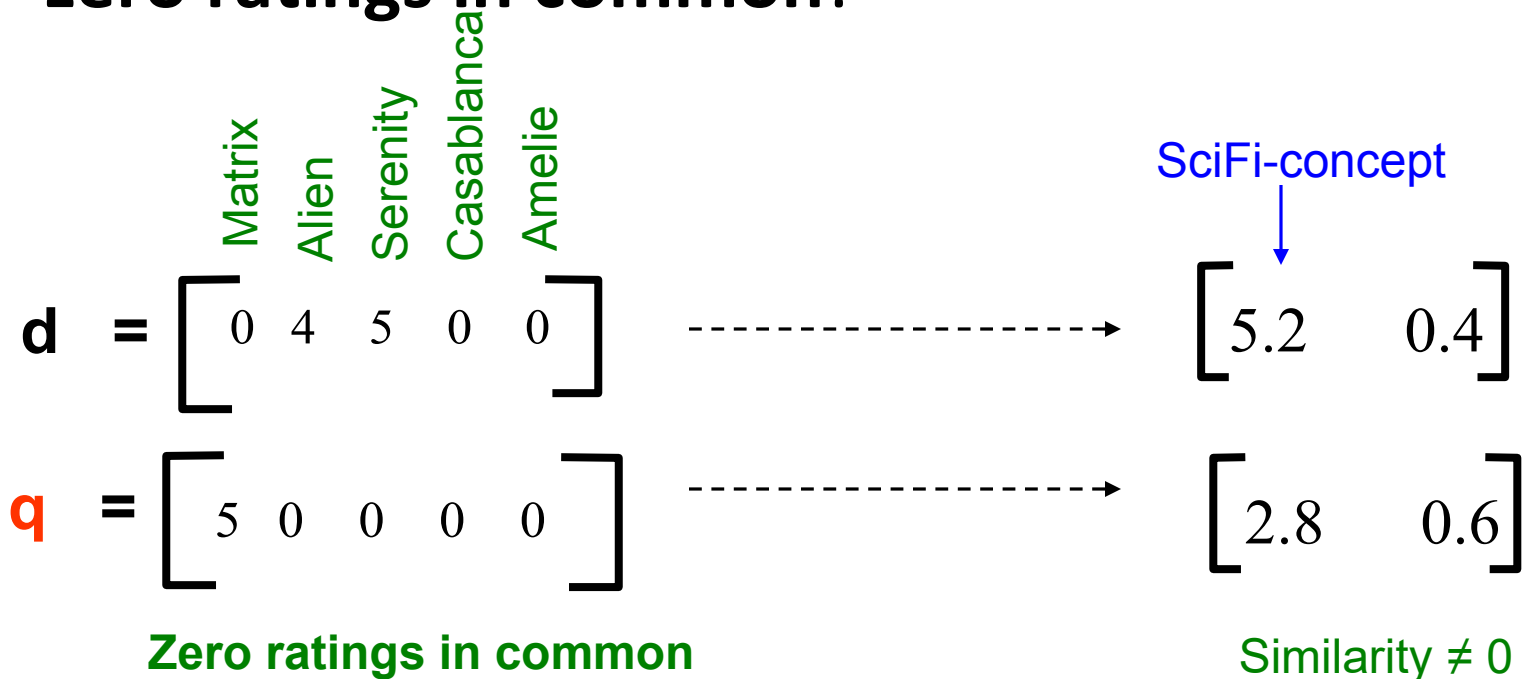
movie-to-concept similarities (V)

SciFi-concept



Case study: How to query?

- **Observation:** User d that rated (*'Alien'*, *'Serenity'*) will be **similar** to user q that rated (*'Matrix'*), although d and q have **zero ratings in common!**





SVD: Pros and Cons

- + **Optimal low-rank approximation**
in terms of Frobenius norm (or Euclidean distance)

Frobenius norm:

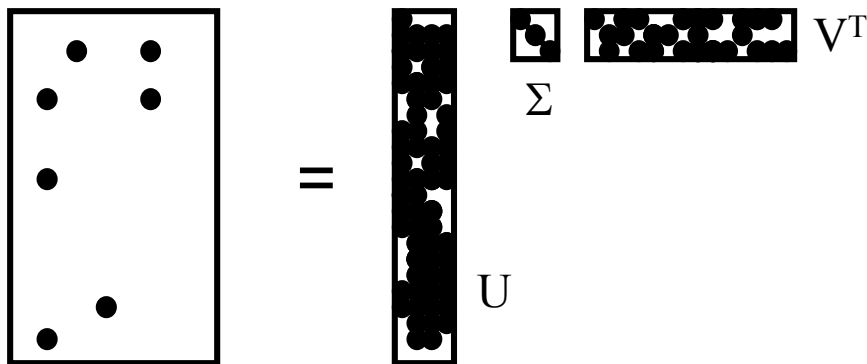
$$\|X\|_F = \sqrt{\sum_{ij} X_{ij}^2}$$

- **Interpretability problem:**

- A singular vector specifies a linear combination of all input columns or rows

- **Lack of sparsity:**

- Singular vectors are **dense!**





Outline

☒ SVD Case Studies

 ☐ **CUR Decomposition**



CUR Decomposition

Frobenius norm: $\|X\|_F = \sqrt{\sum_{ij} X_{ij}^2}$

- Goal: Express A as a product of matrices C, U, R

Make $\|A - C \cdot U \cdot R\|_F$ small

- “Constraints” on C and R :

$$\left(\begin{array}{|c|} \hline \text{Red} \\ \hline \end{array} \quad \begin{array}{|c|} \hline \text{Brown} \\ \hline \end{array} \quad \begin{array}{|c|} \hline \text{Green} \\ \hline \end{array} \right) \approx \left(\begin{array}{|c|c|c|c|c|c|} \hline \text{Red} & \text{Red} & \text{Red} & \text{Brown} & \text{Green} & \text{Green} \\ \hline \end{array} \right) \cdot \left(\begin{array}{c} U \end{array} \right) \cdot \left(\begin{array}{c} R \end{array} \right)$$

$A \qquad C \qquad U \qquad R$



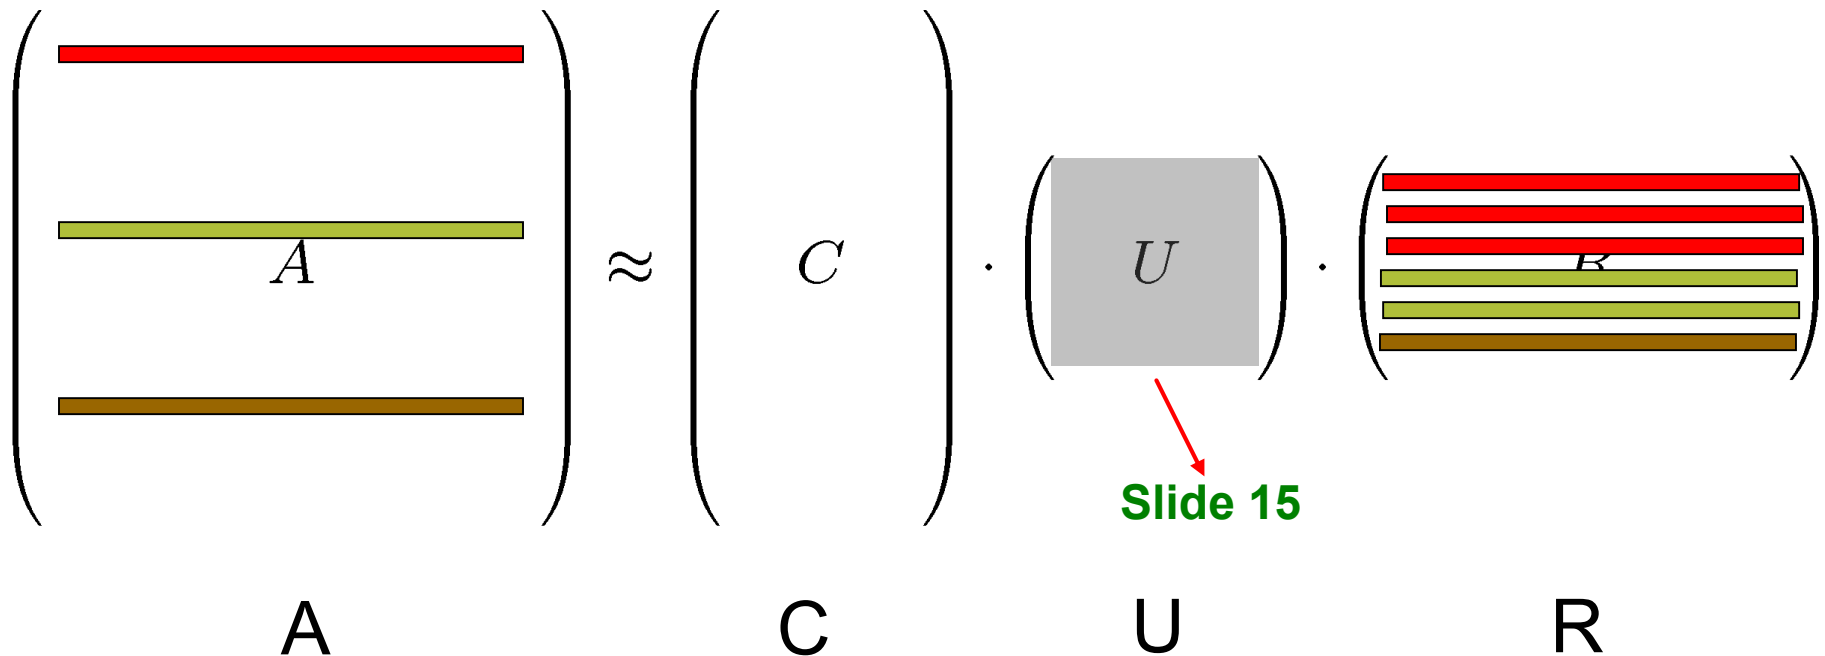
CUR Decomposition

Frobenius norm: $\|X\|_F = \sqrt{\sum_{ij} X_{ij}^2}$

- Goal: Express A as a product of matrices C,U,R

Make $\|A - C \cdot U \cdot R\|_F$ small

- “Constraints” on C and R:





CUR: How it Works

■ Sampling columns (similarly for rows):

Input: matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$, sample size c

Output: $\mathbf{C}_d \in \mathbb{R}^{m \times c}$

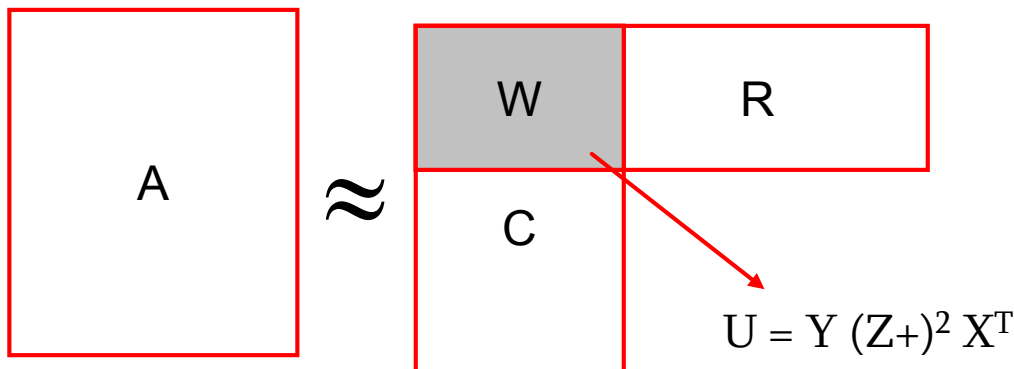
1. for $x = 1 : n$ [column distribution]
2. $P(x) = \sum_i \mathbf{A}(i, x)^2 / \sum_{i,j} \mathbf{A}(i, j)^2$
3. for $i = 1 : c$ [sample columns]
4. Pick $j \in 1 : n$ based on distribution $P(x)$
5. Compute $\mathbf{C}_d(:, i) = \mathbf{A}(:, j) / \sqrt{cP(j)}$

Note this is a randomized algorithm; same column can be sampled more than once



Computing U

- Let **W** be the “intersection” of sampled columns **C** and rows **R**
 - Let SVD of $\mathbf{W} = \mathbf{X} \mathbf{Z} \mathbf{Y}^T$
- **Then: $\mathbf{U} = \mathbf{Y} (\mathbf{Z}^+)^2 \mathbf{X}^T$**
 - \mathbf{Z}^+ : **reciprocals of non-zero singular values: $Z^+_{ii} = 1 / Z_{ii}$**
 - \mathbf{Z}^+ is called “**pseudoinverse**” of \mathbf{Z}





CUR: Provably good approx. to SVD

- If we carefully choose # of columns and rows,

$$\|A - CUR\|_F \leq (2 + \epsilon) \|A - A_k\|_F$$

CUR error SVD error

- with probability 98%

In practice:

Pick $4k$ cols/rows

for a “rank- k ” approximation



CUR: Pros & Cons

+ Easy interpretation

- Since the basis vectors are actual columns and rows

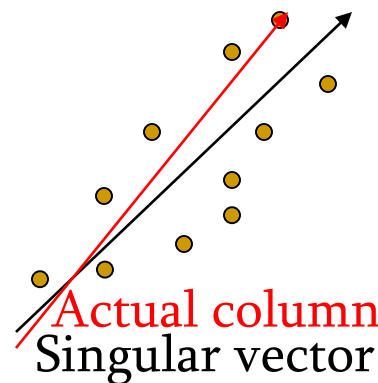
+ Sparse basis

- Since the basis vectors are actual columns and rows

- Not the 'optimal' solution for minimizing error

- Duplicate columns and rows

- Columns of large norms will be sampled many times
 - CMD method (Sun et al., 2007) solves the problem





SVD vs. CUR

sparse and small

SVD: $A = U \Sigma V^T$

Huge but sparse Big and dense

dense but small

CUR: $A = C U R$

Huge but sparse Big but sparse



SVD vs. CUR: Simple Experiment

■ DBLP bibliographic data

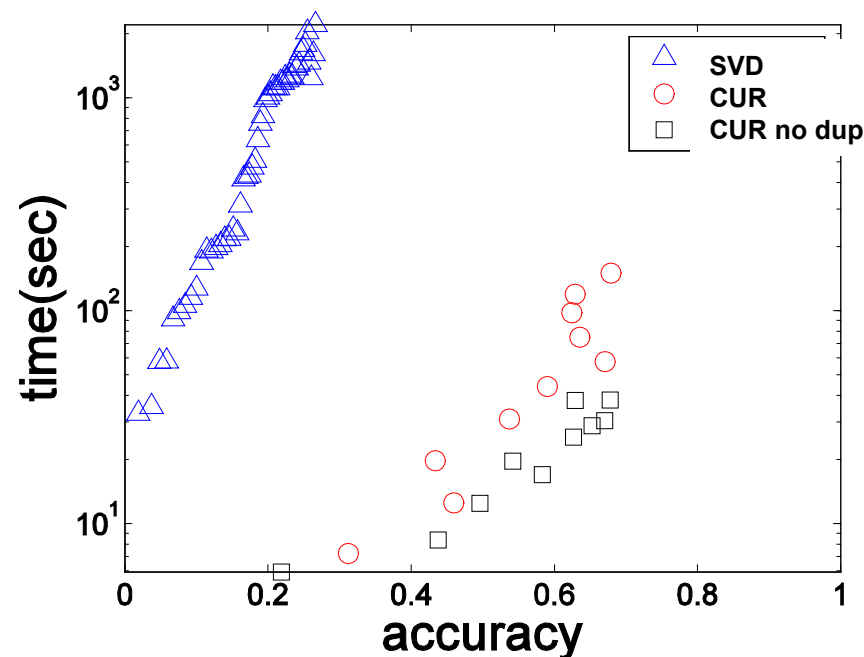
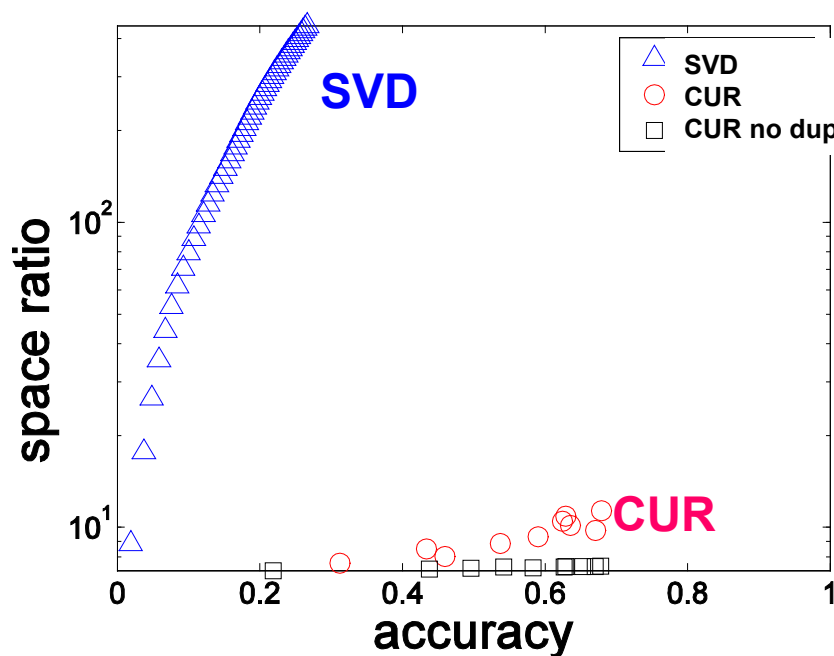
- ❑ Author-to-conference big sparse matrix
- ❑ A_{ij} : Number of papers published by author i at conference j
- ❑ 428K authors (rows), 3659 conferences (columns)
 - Very sparse

■ Want to reduce dimensionality

- ❑ How much time does it take?
- ❑ What is the reconstruction error?
- ❑ How much space do we need?



Results: DBLP- big sparse matrix



■ Accuracy:

- $1 - \text{relative sum squared errors}$

■ Space ratio:

SVD는 hyper parameter를 정해줘야 할 필요가 없고 error bound가 확실하므로 쓰기 훨씬 편하다.

- $\text{\#output matrix entries} / \text{\#input matrix entries}$

■ CPU time

Sun, Faloutsos: *Less is More: Compact Matrix Decomposition for Large Sparse Graphs*, SDM '07.



Further Reading: CUR

- Drineas et al., *Fast Monte Carlo Algorithms for Matrices III: Computing a Compressed Approximate Matrix Decomposition*, SIAM Journal on Computing, 2006.
- J. Sun, Y. Xie, H. Zhang, C. Faloutsos: *Less is More: Compact Matrix Decomposition for Large Sparse Graphs*, SDM 2007
- *Intra- and interpopulation genotype reconstruction from tagging SNPs*, P. Paschou, M. W. Mahoney, A. Javed, J. R. Kidd, A. J. Pakstis, S. Gu, K. K. Kidd, and P. Drineas, Genome Research, 17(1), 96-107 (2007)
- *Tensor-CUR Decompositions For Tensor-Based Data*, M. W. Mahoney, M. Maggioni, and P. Drineas, Proc. 12-th Annual SIGKDD, 327-336 (2006)

Conclusion



Data contains value and knowledge



What did we learn?

- We learned to **mine different types of data:**

- ❑ Data is high dimensional
- ❑ Data is a graph
- ❑ Data is infinite/never-ending

- We learned to **use different models of computation:**

- ❑ MapReduce
- ❑ Streams and online algorithms
- ❑ Single machine in-memory



What did we learn?

- We learned to **solve real-world problems:**
 - ❑ Recommender systems
 - ❑ Market basket analysis
 - ❑ Spam detection
 - ❑ Duplicate document detection
 - ❑ Ranking in graphs
 - ❑ Community detection
 - ❑ Increasing the revenue of search engines
 - ❑ Trending topics in social network news feeds



What did we learn?

■ We learned **various “tools”**:

- ❑ Linear algebra (SVD, Rec. Sys., Communities)
- ❑ Pruning technique (frequent itemsets)
- ❑ Hashing (LSH, Bloom filters)
- ❑ Link Analysis (PageRank, HITS, Random Walk)
- ❑ Graph Algorithms (betweenness, BFS)
- ❑ Clustering (k-means, CURE)
- ❑ Sampling (reservoir sampling)
- ❑ Greedy Algorithm (Balance)



Final Remark

- This course covered extensive materials; thanks for all the hard works (quizzes, homeworks, exams, questions, etc.)
- The knowledge you learned from this course would be helpful in many ways
 - E.g., which subjects to study more; understanding and evaluating other people's work
- Takeaway points
 - Summarize each topic using few sentences so that you can really use them (e.g., Bloom Filter)
 - 1) What is the problem to solve? 2) What is the main idea? 3) Result?
 - Learn the way how each problem is solved
 - To apply the similar techniques to solve other problems in the future



Thank You!